

# A Brief Overview of Recent Neural Network Methods and Architectures

UCLA CS161

Mohammad Kachuee

Dec 6, 2017


# Overview

- 1 Fully Connected Neural Networks
  - Multi-layer Preceptron (MLP)
  - Limitations of MLP
  - Inspirations from Visual Cortex
- 2 Convolutional Neural Networks (CNNs)
  - Basic Ideas of CNNs
  - Famous CNN Architectures
  - Visualizing CNNs
- 3 Recurrent Neural Networks (RNNs)
  - Handling Sequences as Input
  - RNN Applications
  - Basic Ideas of RNNs
- 4 TensorFlow Tutorial

- 1 **Fully Connected Neural Networks**
  - **Multi-layer Preceptron (MLP)**
  - **Limitations of MLP**
  - **Inspirations from Visual Cortex**
- 2 **Convolutional Neural Networks (CNNs)**
  - Basic Ideas of CNNs
  - Famous CNN Architectures
  - Visualizing CNNs
- 3 **Recurrent Neural Networks (RNNs)**
  - Handling Sequences as Input
  - RNN Applications
  - Basic Ideas of RNNs
- 4 **TensorFlow Tutorial**

# Multi-layer Preceptron (MLP) (1/2)

$$h_{i+1} = \sigma(h_i W_i + b_i)$$



# Multi-layer Preceptron (MLP) (2/2)

- In a MLP, we have a few fully connected layers that are stacked with each other.
- There are three layer types: input layer, hidden layer, and output layer.
- Each layer takes the activation values of the previous layer as input, applies a linear transformation to it, and crates output activations by applying a non-linearity.
- The universal approximation theorem: theoretically, we can model any function using a neural network with a single hidden layer.

# Limitations of MLP

- MLP discards any structure in the input space and creates a fully connected many to many mapping.

# Limitations of MLP

- MLP discards any structure in the input space and creates a fully connected many to many mapping.
- Each layer in MLP has  $O(mn)$  weights ( $m$  input neurons and  $n$  output neurons).

# Limitations of MLP

- MLP discards any structure in the input space and creates a fully connected many to many mapping.
- Each layer in MLP has  $O(mn)$  weights ( $m$  input neurons and  $n$  output neurons).
- Having many neurons in each layer and many layers makes the model very complex and hard to train.



# Limitations of MLP

- MLP discards any structure in the input space and creates a fully connected many to many mapping.
- Each layer in MLP has  $O(mn)$  weights ( $m$  input neurons and  $n$  output neurons).
- Having many neurons in each layer and many layers makes the model very complex and hard to train.
- Practically, training networks with more than a few layers is not easy.




# Limitations of MLP

- MLP discards any structure in the input space and creates a fully connected many to many mapping.
- Each layer in MLP has  $O(mn)$  weights ( $m$  input neurons and  $n$  output neurons).
- Having many neurons in each layer and many layers makes the model very complex and hard to train.
- Practically, training networks with more than a few layers is not easy.
- Finding the optimal number of layers and neurons inside each one is an open problem.

# Limitations of MLP

- MLP discards any structure in the input space and creates a fully connected many to many mapping.
- Each layer in MLP has  $O(mn)$  weights ( $m$  input neurons and  $n$  output neurons).
- Having many neurons in each layer and many layers makes the model very complex and hard to train.
- Practically, training networks with more than a few layers is not easy.
- Finding the optimal number of layers and neurons inside each one is an open problem.
- The universal approximation theorem is not helping us that much because a huge number of hidden neurons needed to approximate a function using a shallow network.


# Hubel and Wiesel Experiments



# Hierarchical Organization of Visual Cortex

## Hubel & Weisel

topographical mapping




## featural hierarchy

hyper-complex cells

complex cells

simple cells



high level



mid level




low level

- 1 Fully Connected Neural Networks
  - Multi-layer Perceptron (MLP)
  - Limitations of MLP
  - Inspirations from Visual Cortex
- 2 **Convolutional Neural Networks (CNNs)**
  - **Basic Ideas of CNNs**
  - **Famous CNN Architectures**
  - **Visualizing CNNs**
- 3 Recurrent Neural Networks (RNNs)
  - Handling Sequences as Input
  - RNN Applications
  - Basic Ideas of RNNs
- 4 TensorFlow Tutorial


# The Era of Deep Representations (1/2)

- Image recognition is a machine learning application in which the machine should predict objects that are present in the image.
- Before 2012, the major method in image recognition was using hand engineered features and classic machine learning classifiers.




# The Era of Deep Representations (1/2)

- Image recognition is a machine learning application in which the machine should predict objects that are present in the image.
- Before 2012, the major method in image recognition was using hand engineered features and classic machine learning classifiers.




- Around 2012, three major factors changed the future landscape of recognition tasks from engineered features and shallow architectures to using automated feature extraction through deep neural networks.
  - Significantly bigger datasets such as ImageNet (10,000,000 labeled images depicting 10,000+ object categories)
  - More powerful computers and GPU accelerated computation.
  - Development of techniques and tricks for training deeper networks.





# The Era of Deep Representations (2/2)



Model	Top-1	Top-5
<i>Sparse coding [2]</i>	47.1%	28.2%
<i>SIFT + FVs [24]</i>	45.7%	25.7%
CNN	<b>37.5%</b>	<b>17.0%</b>

Table 1: Comparison of results on ILSVRC-2010 test set. In *italics* are best results achieved by others.

# Natural Images as Input

- Natural images are smooth.



# Natural Images as Input

- Natural images are smooth.
- Nearby pixels are coherent and highly correlated with each other.



# Natural Images as Input

- Natural images are smooth.
- Nearby pixels are coherent and highly correlated with each other.
- An object inside an image can be at different coordinates.



# Natural Images as Input

- Natural images are smooth.
- Nearby pixels are coherent and highly correlated with each other.
- An object inside an image can be at different coordinates.
- An object inside an image can be at different scales or distances from the view-point.



# Natural Images as Input

- Natural images are smooth.
- Nearby pixels are coherent and highly correlated with each other.
- An object inside an image can be at different coordinates.
- An object inside an image can be at different scales or distances from the view-point.
- In real-world, we have lighting and noise affecting the image.



# Natural Images as Input

- Natural images are smooth.
- Nearby pixels are coherent and highly correlated with each other.
- An object inside an image can be at different coordinates.
- An object inside an image can be at different scales or distances from the view-point.
- In real-world, we have lighting and noise affecting the image.

Is it possible to take the advantage of image characteristics in designing the network?



# Basic Ideas of CNNs: Intuition

- The same object may be located anywhere inside an image.



# Basic Ideas of CNNs: Intuition

- The same object may be located anywhere inside an image.
- Image patch characteristics are almost similar across different scenes and views.

# Basic Ideas of CNNs: Intuition

- The same object may be located anywhere inside an image.
- Image patch characteristics are almost similar across different scenes and views.
- Use the same feature detector at different input locations.

# Basic Ideas of CNNs: Intuition

- The same object may be located anywhere inside an image.
- Image patch characteristics are almost similar across different scenes and views.
- Use the same feature detector at different input locations.
- In other words, share network weights to prevent learning redundant mappings and reuse feature detectors.

# Basic Ideas of CNNs: Intuition


- The same object may be located anywhere inside an image.
- Image patch characteristics are almost similar across different scenes and views.
- Use the same feature detector at different input locations.
- In other words, share network weights to prevent learning redundant mappings and reuse feature detectors.
- It can be easily modeled as the 2D convolution operation.

# Basic Ideas of CNNs: Intuition

- The same object may be located anywhere inside an image.
- Image patch characteristics are almost similar across different scenes and views.
- Use the same feature detector at different input locations.
- In other words, share network weights to prevent learning redundant mappings and reuse feature detectors.
- It can be easily modeled as the 2D convolution operation.
- It is, basically, learning a set of filters and applying them at different image locations.


# Basic Ideas of CNNs: Conv Layers (1/3)

Applying a convolution filter on an image:




## Basic Ideas of CNNs: Conv Layers (2/3)

Applying two different convolution filters on an image:



# Basic Ideas of CNNs: Conv Layers (3/3)

Applying a set of different convolution filters on an image:



Fei-Fei Li, Andrej Karpathy, and Justin Johnson



## Benefits of CNNs: Reducing model complexity

- In the example above, we mapped from a  $32 \times 32 \times 3$  input to a  $28 \times 28 \times 1$  output using a  $5 \times 5 \times 3$  filter.

# Benefits of CNNs: Reducing model complexity

- In the example above, we mapped from a  $32 \times 32 \times 3$  input to a  $28 \times 28 \times 1$  output using a  $5 \times 5 \times 3$  filter.
- If we had a fully connected mapping, similar to what we had in MLP, we needed a  $(32 \times 32 \times 3) \times (28 \times 28 \times 1)$  weight matrix.

# Benefits of CNNs: Reducing model complexity

- In the example above, we mapped from a  $32 \times 32 \times 3$  input to a  $28 \times 28 \times 1$  output using a  $5 \times 5 \times 3$  filter.
- If we had a fully connected mapping, similar to what we had in MLP, we needed a  $(32 \times 32 \times 3) \times (28 \times 28 \times 1)$  weight matrix.
- It means using 75 weights versus 2408448!

# Benefits of CNNs: Reducing model complexity

- In the example above, we mapped from a  $32 \times 32 \times 3$  input to a  $28 \times 28 \times 1$  output using a  $5 \times 5 \times 3$  filter.
- If we had a fully connected mapping, similar to what we had in MLP, we needed a  $(32 \times 32 \times 3) \times (28 \times 28 \times 1)$  weight matrix.
- It means using 75 weights versus 2408448!
- Usually there are more than one convolution filter.

# Benefits of CNNs: Reducing model complexity

- In the example above, we mapped from a  $32 \times 32 \times 3$  input to a  $28 \times 28 \times 1$  output using a  $5 \times 5 \times 3$  filter.
- If we had a fully connected mapping, similar to what we had in MLP, we needed a  $(32 \times 32 \times 3) \times (28 \times 28 \times 1)$  weight matrix.
- It means using 75 weights versus 2408448!
- Usually there are more than one convolution filter.
- However, the overall complexity is an order of magnitude lower!

# Benefits of CNNs: Reducing model complexity

- In the example above, we mapped from a  $32 \times 32 \times 3$  input to a  $28 \times 28 \times 1$  output using a  $5 \times 5 \times 3$  filter.
- If we had a fully connected mapping, similar to what we had in MLP, we needed a  $(32 \times 32 \times 3) \times (28 \times 28 \times 1)$  weight matrix.
- It means using 75 weights versus 2408448!
- Usually there are more than one convolution filter.
- However, the overall complexity is an order of magnitude lower!
- Lower complexity means ability to fit models faster, easier, and using fewer training samples.

# Benefits of CNNs: Fast GPU Training

- Modern GPUs are optimized for fast and efficient convolution.




Figure: Nvidia GTX 1080Ti (3584 Cores, 11GB Memory). Source: [nvidia.com](https://www.nvidia.com)

# Benefits of CNNs: Fast GPU Training

- Modern GPUs are optimized for fast and efficient convolution.
- A modern GPU has about a few thousands of cores and tens of Gigabytes Ram that can work in parallel.




Figure: Nvidia GTX 1080Ti (3584 Cores, 11GB Memory). Source: [nvidia.com](https://www.nvidia.com)



# Benefits of CNNs: Fast GPU Training

- Modern GPUs are optimized for fast and efficient convolution.
- A modern GPU has about a few thousands of cores and tens of Gigabytes Ram that can work in parallel.
- Many GPUs can be coupled with each other, creating a GPU farm.





Figure: Nvidia GTX 1080Ti (3584 Cores, 11GB Memory). Source: [nvidia.com](https://www.nvidia.com)

# Famous CNN Architectures: LeNet 1998




LeCun, Yann, et al. "Gradient-based learning applied to document recognition." Proceedings of the IEEE 86.11 (1998): 2278-2324.



# Famous CNN Architectures: ImageNet 2012 (AlexNet) (2/2)

Model	Top-1	Top-5
<i>Sparse coding [2]</i>	47.1%	28.2%
<i>SIFT + FVs [24]</i>	45.7%	25.7%
CNN	<b>37.5%</b>	<b>17.0%</b>


Table 1: Comparison of results on ILSVRC-2010 test set. In *italics* are best results achieved by others.



Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." *Advances in neural information processing systems*. 2012.

# Famous CNN Architectures: VGG (1/2)

About 20 weight layers and 130 million parameters.



Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." arXiv preprint arXiv:1409.1556 (2014).

# Famous CNN Architectures: VGG (2/2)

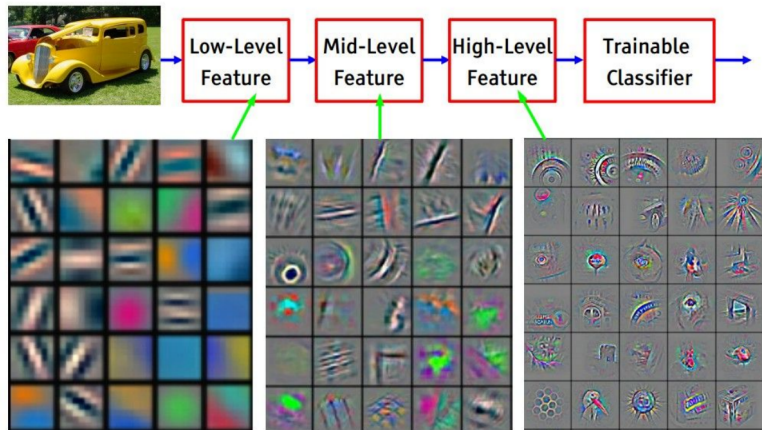
Table 7: **Comparison with the state of the art in ILSVRC classification.** Our method is denoted as “VGG”. Only the results obtained without outside training data are reported.

Method	top-1 val. error (%)	top-5 val. error (%)	top-5 test error (%)
VGG (2 nets, multi-crop & dense eval.)	<b>23.7</b>	<b>6.8</b>	<b>6.8</b>
VGG (1 net, multi-crop & dense eval.)	24.4	7.1	7.0
VGG (ILSVRC submission, 7 nets, dense eval.)	24.7	7.5	7.3
GoogLeNet (Szegedy et al., 2014) (1 net)	-	7.9	
GoogLeNet (Szegedy et al., 2014) (7 nets)	-	<b>6.7</b>	
MSRA (He et al., 2014) (11 nets)	-	-	8.1
MSRA (He et al., 2014) (1 net)	27.9	9.1	9.1
Clarifai (Russakovsky et al., 2014) (multiple nets)	-	-	11.7
Clarifai (Russakovsky et al., 2014) (1 net)	-	-	12.5
Zeiler & Fergus (Zeiler & Fergus, 2013) (6 nets)	36.0	14.7	14.8
Zeiler & Fergus (Zeiler & Fergus, 2013) (1 net)	37.5	16.0	16.1
OverFeat (Sermanet et al., 2014) (7 nets)	34.0	13.2	13.6
OverFeat (Sermanet et al., 2014) (1 net)	35.7	14.2	-
Krizhevsky et al. (Krizhevsky et al., 2012) (5 nets)	38.1	16.4	16.4
Krizhevsky et al. (Krizhevsky et al., 2012) (1 net)	40.7	18.2	-

---

Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." arXiv preprint arXiv:1409.1556 (2014).

# Visualizing CNNs



Zeiler, Matthew D., and Rob Fergus. "Visualizing and understanding convolutional networks." European conference on computer vision. Springer, Cham, 2014.

# Outline

- 1 Fully Connected Neural Networks
  - Multi-layer Perceptron (MLP)
  - Limitations of MLP
  - Inspirations from Visual Cortex
- 2 Convolutional Neural Networks (CNNs)
  - Basic Ideas of CNNs
  - Famous CNN Architectures
  - Visualizing CNNs
- 3 Recurrent Neural Networks (RNNs)
  - Handling Sequences as Input
  - RNN Applications
  - Basic Ideas of RNNs
- 4 TensorFlow Tutorial



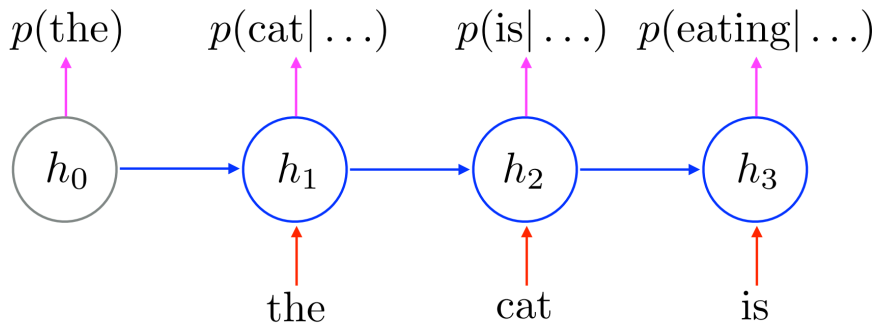
# Handling Sequences as Input

- We have a wide range of applications in which:
  - We are dealing with a sequence of inputs (not a single feature vector!)
  - The sequence length is not necessarily fixed.
  - We may need to have multiple output values at different time instances.

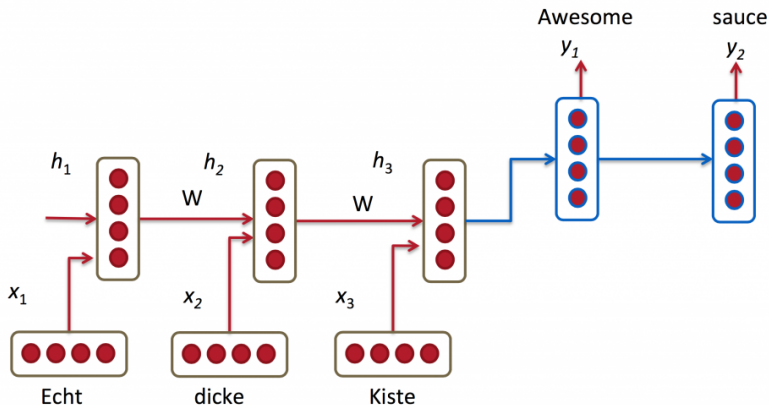
# Handling Sequences as Input

- We have a wide range of applications in which:
  - We are dealing with a sequence of inputs (not a single feature vector!)
  - The sequence length is not necessarily fixed.
  - We may need to have multiple output values at different time instances.
- For instance:
  - Predictive analysis in time series. e.g., predicting the expected future value of the stocks for a company.
  - Language modeling. e.g., designing an auto-complete engine which can predict and suggest the next word given the current incomplete sentence.
  - Sequence to sequence modeling. e.g., translating a speech from English to French in real-time.

# RNN Applications: Language Modeling



# RNN Applications: Sentence Translation



# Basic Ideas of RNNs

- RNNs are capable of having sequences as input or output.

# Basic Ideas of RNNs

- RNNs are capable of having sequences as input or output.
- Recurrent connections in neural networks can be used to create a notion of memory.

# Basic Ideas of RNNs

- RNNs are capable of having sequences as input or output.
- Recurrent connections in neural networks can be used to create a notion of memory.
- RNNs can be trained by using the same general techniques we had before (unrolling technique).

# Basic Ideas of RNNs

- RNNs are capable of having sequences as input or output.
- Recurrent connections in neural networks can be used to create a notion of memory.
- RNNs can be trained by using the same general techniques we had before (unrolling technique).
- A more detailed discussion of RNNs is out of the scope of this lecture.




# Outline

- 1 Fully Connected Neural Networks
  - Multi-layer Perceptron (MLP)
  - Limitations of MLP
  - Inspirations from Visual Cortex
- 2 Convolutional Neural Networks (CNNs)
  - Basic Ideas of CNNs
  - Famous CNN Architectures
  - Visualizing CNNs
- 3 Recurrent Neural Networks (RNNs)
  - Handling Sequences as Input
  - RNN Applications
  - Basic Ideas of RNNs
- 4 **TensorFlow Tutorial**

# TensorFlow: Introduction (1/2)


- Consider the computation of forward path in a simple one layer neural network:

$$Y = \tanh(WX + b)$$




# TensorFlow: Introduction (1/2)

- Consider the computation of forward path in a simple one layer neural network:  
 $Y = \tanh(WX + b)$
- A data flow graph is a computational graph that explains the processing that is happening to the flow of data.




# TensorFlow: Introduction (1/2)

- Consider the computation of forward path in a simple one layer neural network:  
 $Y = \tanh(WX + b)$
- A data flow graph is a computational graph that explains the processing that is happening to the flow of data.
- Data flow model is in contrast to the control flow model.



# TensorFlow: Introduction (1/2)


- Consider the computation of forward path in a simple one layer neural network:  
 $Y = \tanh(WX + b)$
- A data flow graph is a computational graph that explains the processing that is happening to the flow of data.
- Data flow model is in contrast to the control flow model.
- TensorFlow is a computational library by Google that is designed for fast and scalable computing based on data flow graphs.



- In a typical computational library, we call functions on data sequentially.
  - It adds the overhead of function call, argument pass, and return to the actual processing time.
  - It prevents optimizing in higher-levels. E.g., multiplication result of two huge matrices is multiplied by a vector.

## TensorFlow: Introduction (2/2)

- In a typical computational library, we call functions on data sequentially.
  - It adds the overhead of function call, argument pass, and return to the actual processing time.
  - It prevents optimizing in higher-levels. E.g., multiplication result of two huge matrices is multiplied by a vector.
- In TensorFlow, we create a complete computational graph. Then we feed our data to the graph and every process takes place in background.



# Basics: Variables, Constants, Operations

- We can define variable and constant tensors similar to numpy or any other computational library.



# Basics: Variables, Constants, Operations

- We can define variable and constant tensors similar to numpy or any other computational library.
- We can define operations between tensors to create new tensors.

# Basics: Variables, Constants, Operations

- We can define variable and constant tensors similar to numpy or any other computational library.
- We can define operations between tensors to create new tensors.
- Learning and optimization functions and operations are also can be defined in this fashion.

# Basics: Variables, Constants, Operations

- We can define variable and constant tensors similar to numpy or any other computational library.
- We can define operations between tensors to create new tensors.
- Learning and optimization functions and operations are also can be defined in this fashion.
- Defined variables are by default trainable and need to be initialized before usage.

```
1 # import tensorflow package
2 import tensorflow as tf
3 # create a constant tensor of ones
4 C = tf.ones((2,2))
5 # create a tensor variable initialized with zeros, name it
  weights
6 W = tf.Variable(tf.zeros((2,2)), name="weights")
7 # define a multiplication tensor (Y) which is the result of
  multiplying C by W
8 Y = tf.matmul(C, W, name='op_matmul_CW')
```

# Basics: Placeholders

- Placeholder variables are input ports to the TensorFlow graph.

# Basics: Placeholders

- Placeholder variables are input ports to the TensorFlow graph.
- The Python program can feed data into and get information from the computational graph through placeholders.

# Basics: Placeholders

- Placeholder variables are input ports to the TensorFlow graph.
- The Python program can feed data into and get information from the computational graph through placeholders.
- We define placeholders and use them as regular tensor variables inside our graph code.

# Basics: Placeholders

- Placeholder variables are input ports to the TensorFlow graph.
- The Python program can feed data into and get information from the computational graph through placeholders.
- We define placeholders and use them as regular tensor variables inside our graph code.
- before running the graph, we assign values to the placeholders in our Python code.

```
1
2 # create a place holder of type float and shape 2x2
3 P = tf.placeholder(dtype=tf.float32, shape=(2,2), name='
    ph_input')
4 # create a tensor variable initialized with zeros, name it
    weights
5 W = tf.Variable(tf.zeros((2,2)), name="weights")
6 # define a multiplication tensor (Y) which is the result of
    multiplying P by W
7 Y = tf.matmul(P, W, name='op_matmul_CW')
```

# Basics: Session and Graph

- After creating a graph by chaining TensorFlow operations, we have our computational graph specified.



# Basics: Session and Graph

- After creating a graph by chaining TensorFlow operations, we have our computational graph specified.
- However, before using a graph, we need to create a session in which the graph evaluation will take place.

# Basics: Session and Graph

- After creating a graph by chaining TensorFlow operations, we have our computational graph specified.
- However, before using a graph, we need to create a session in which the graph evaluation will take place.
- In TensorFlow we do this by creating a session object and running it while feeding the placeholders.

```
1 # the next three lines are similar to the previous example
2 P = tf.placeholder(dtype=tf.float32, shape=(2,2), name='
    ph_input')
3 W = tf.Variable(tf.zeros((2,2)), name="weights")
4 Y = tf.matmul(P, W, name='op_matmul_CW')
5 # create a session
6 sess = tf.Session()
7 # create a dummy feed data
8 feed_data = {P: np.ones((2,2), dtype=np.float32)}
9 # run the Y graph using the session and feed the dummy data
10 res = sess.run(Y, feed_dict=feed_data)
11 # print the result
12 print(res)
```

# Training and Optimization

- TensorFlow supports automatic differentiation.
- For training a neural net, we need to:
  - Define forward path.
  - Define a cost function.
  - Create an optimizer.
  - Run the optimization operation to update network weights.

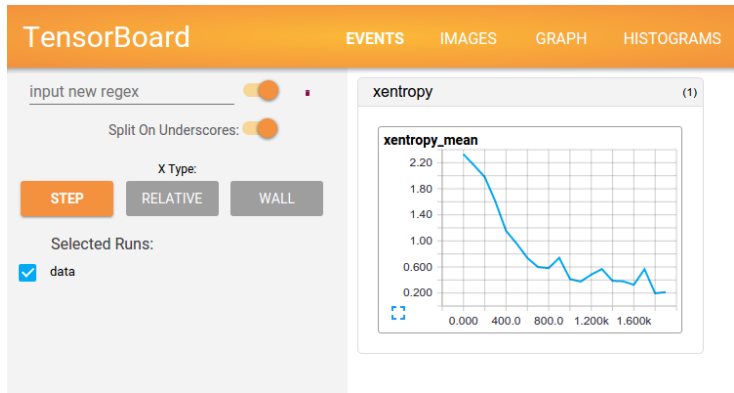
```
1 # define a train step
2 train_step = tf.train.GradientDescentOptimizer(0.5).minimize(
    cross_entropy)
3 # training for 1000 iterations
4 for _ in range(1000):
5     # load a dataset batch
6     batch_xs, batch_ys = mnist.train.next_batch(100)
7     # run the train step
8     sess.run(train_step, feed_dict={x: batch_xs, y_: batch_ys})
```

# Linear Regression using TensorFlow

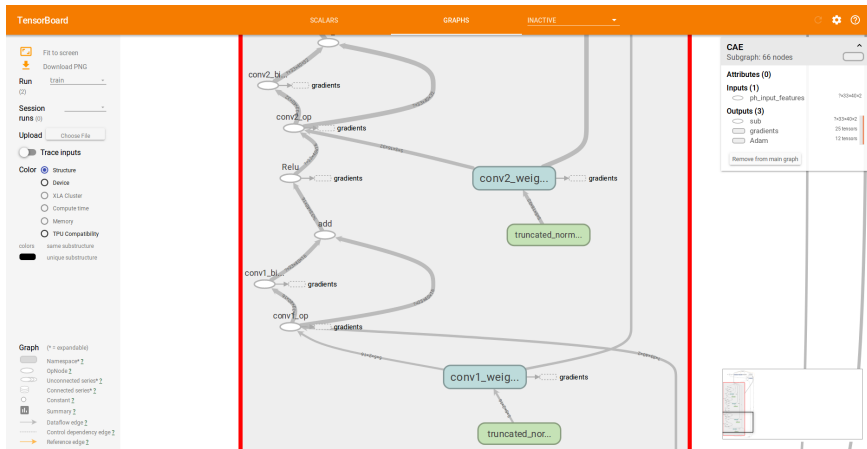
```
1 # define placeholders
2 X = tf.placeholder(dtype=tf.float32, shape=(None,1), name='ph_input')
3 Y = tf.placeholder(dtype=tf.float32, shape=(None,1), name='ph_output')
4 # create variable tensors
5 W = tf.Variable(0.0, dtype=tf.float32, name="weight")
6 b = tf.Variable(0.0, dtype=tf.float32, name="bias")
7 # define a cost function
8 errors = Y - X * W + b
9 cost_mse = tf.reduce_mean(errors ** 2.0)
10 # define a train step
11 train_step = tf.train.GradientDescentOptimizer(0.01).minimize(cost_mse)
12 # training for 1000 iterations
13 for iter_trn in range(1000):
14     # load a dataset batch
15     batch_xs, batch_ys = train_dataset.next_batch(100)
16     # run the train step
17     sess.run(train_step, feed_dict={X: batch_xs, Y: batch_ys})
18     # test the model
19     batch_xs, batch_ys = test_dataset.next_batch(100)
20     cost = sess.run(cost_mse, feed_dict={X: batch_xs, Y: batch_ys})
21     print('Iter: ', iter_trn, ' Test cost: ', cost)
```

# TensorBoard (1/2)

- TensorBoard is a web application by TensorFlow that offers visualization tools for graph visualization, drawing learning curves and histograms, etc.



# TensorBoard (2/2)



# The End